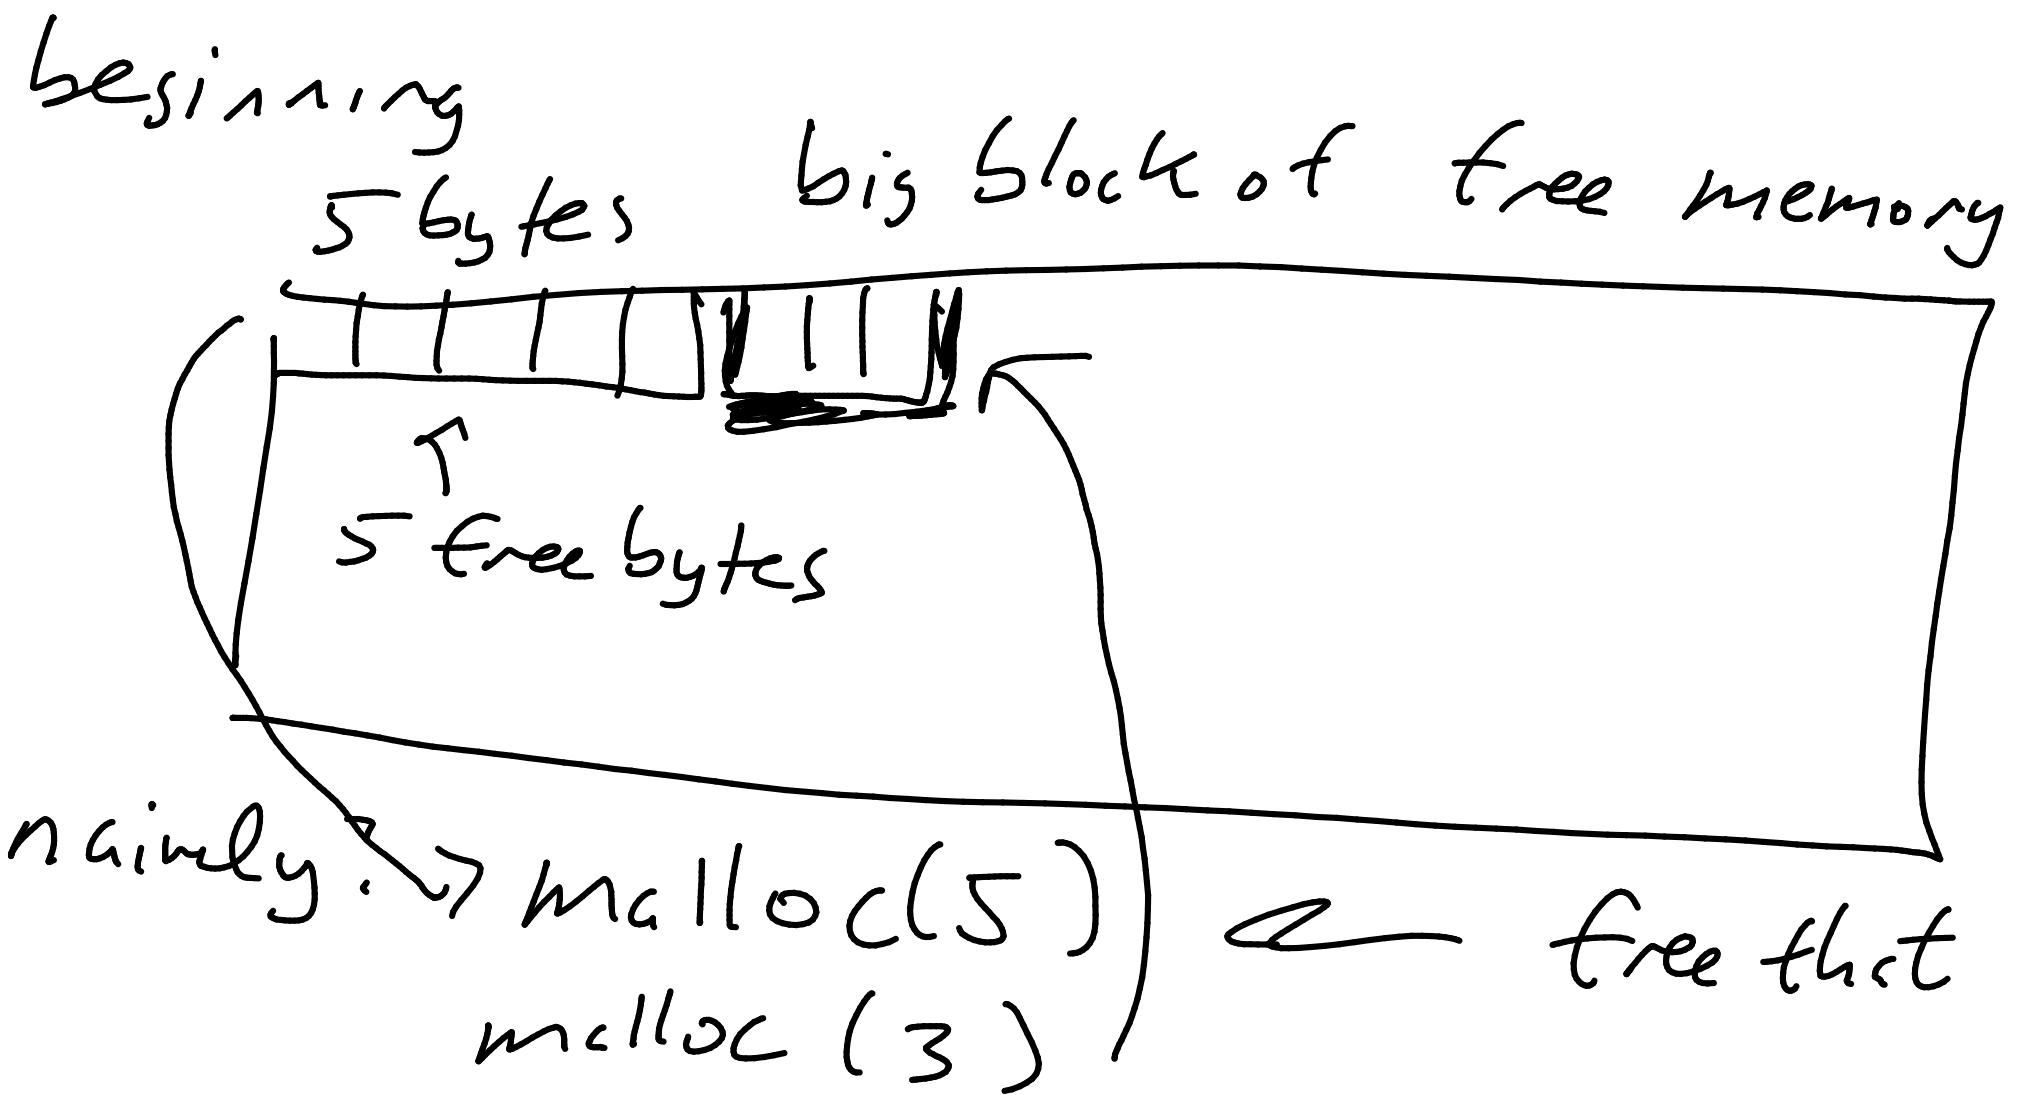


Reminders: today - declare if potentially
self-scheduling

Worksheet on memory allocation

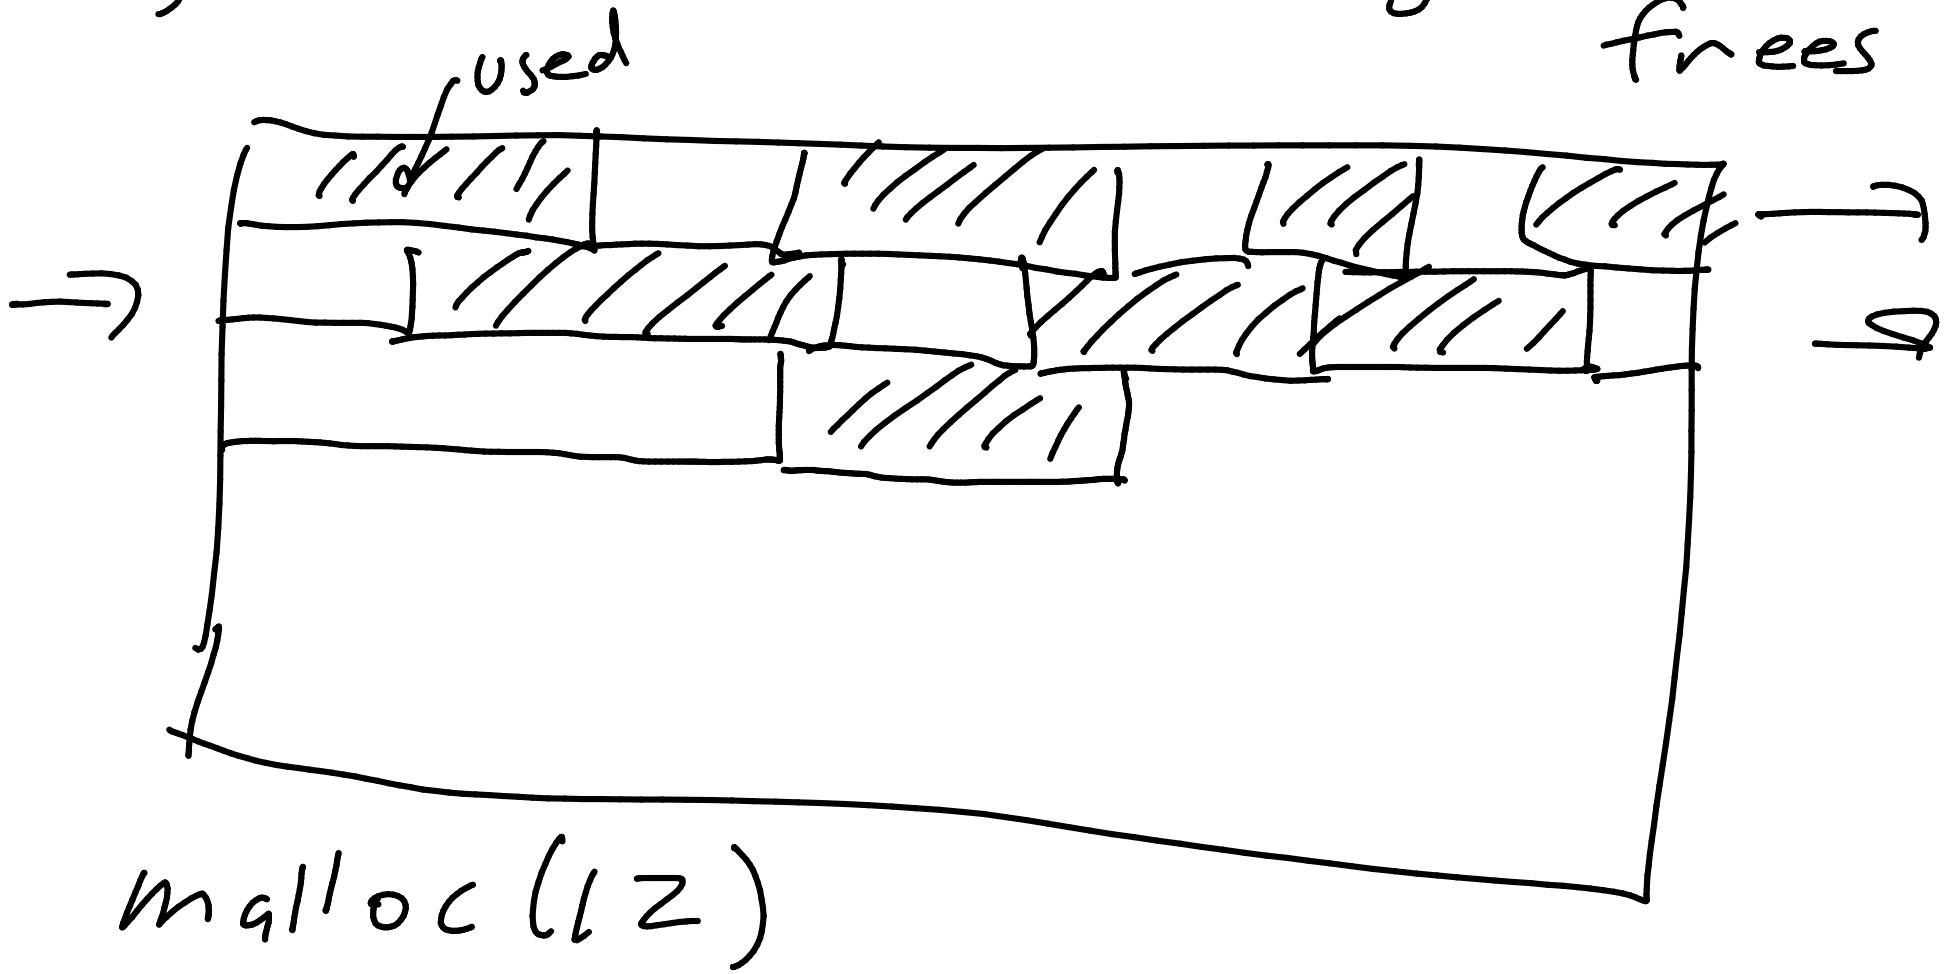
- 5-min-ish
 - w/ groups review it and remember your initial answers
-

How does malloc/memory allocation in general work? How to do it efficiently



`malloc(>)` - search through all
empty regions to see if big
enough - slow if not smart
(binary search is a no go because
data is not sorted)

Fragmentation: after many mallocs and
frees



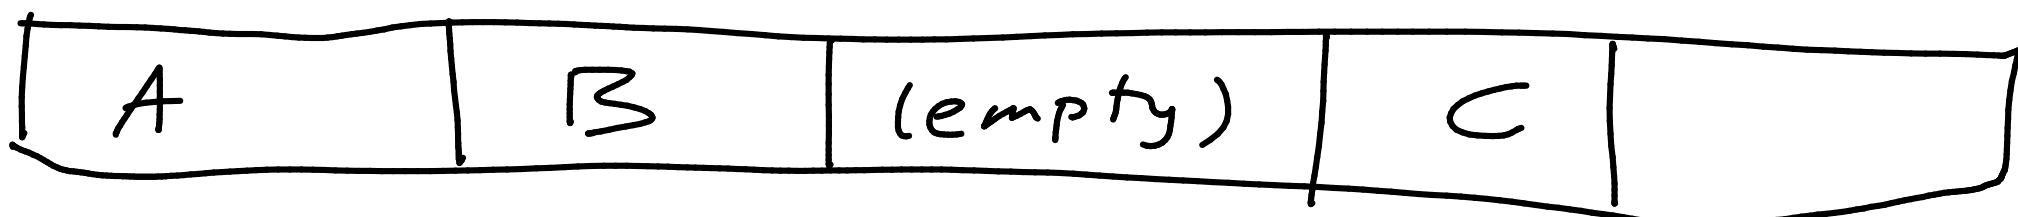
① do I pick the first spot I find
that's at least 12 bytes?

OR

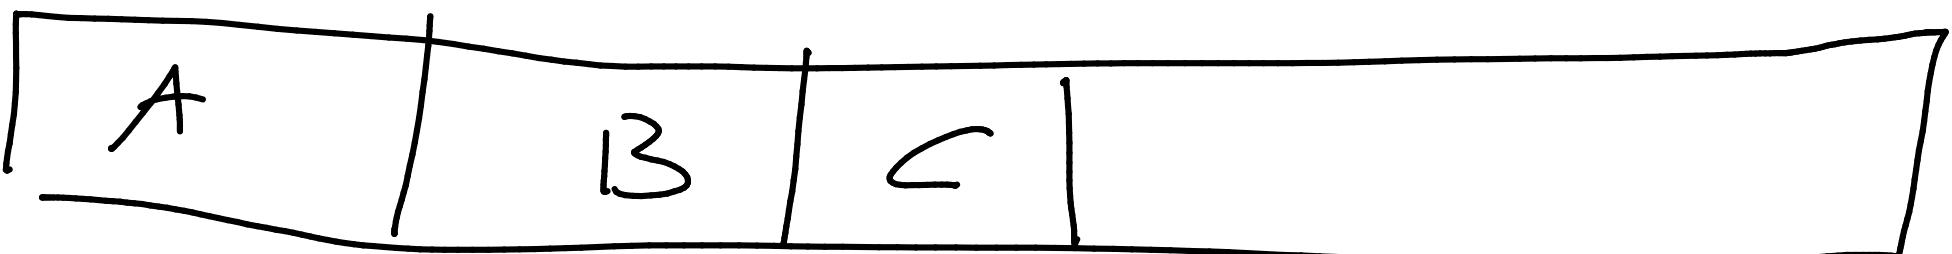
② do I pick the smallest one I
can find that's at least 12

- ① "first fit" - lots of memory left behind, which someone else could use later, but over time these really small and separated
- ② "best fit" - slow (if you're not smart) because need to search everything

Can we defrag memory?



↓ Process ←



Changes pointer
values?

empty

is that ok?

in C, for example

```
int *x = malloc(5);
```

```
printf("%p\n", x);
```

```
printf("%p\n", x);
```

```
printf("%i\n", *x);
```

when you defrag, you have to
change all the variable values

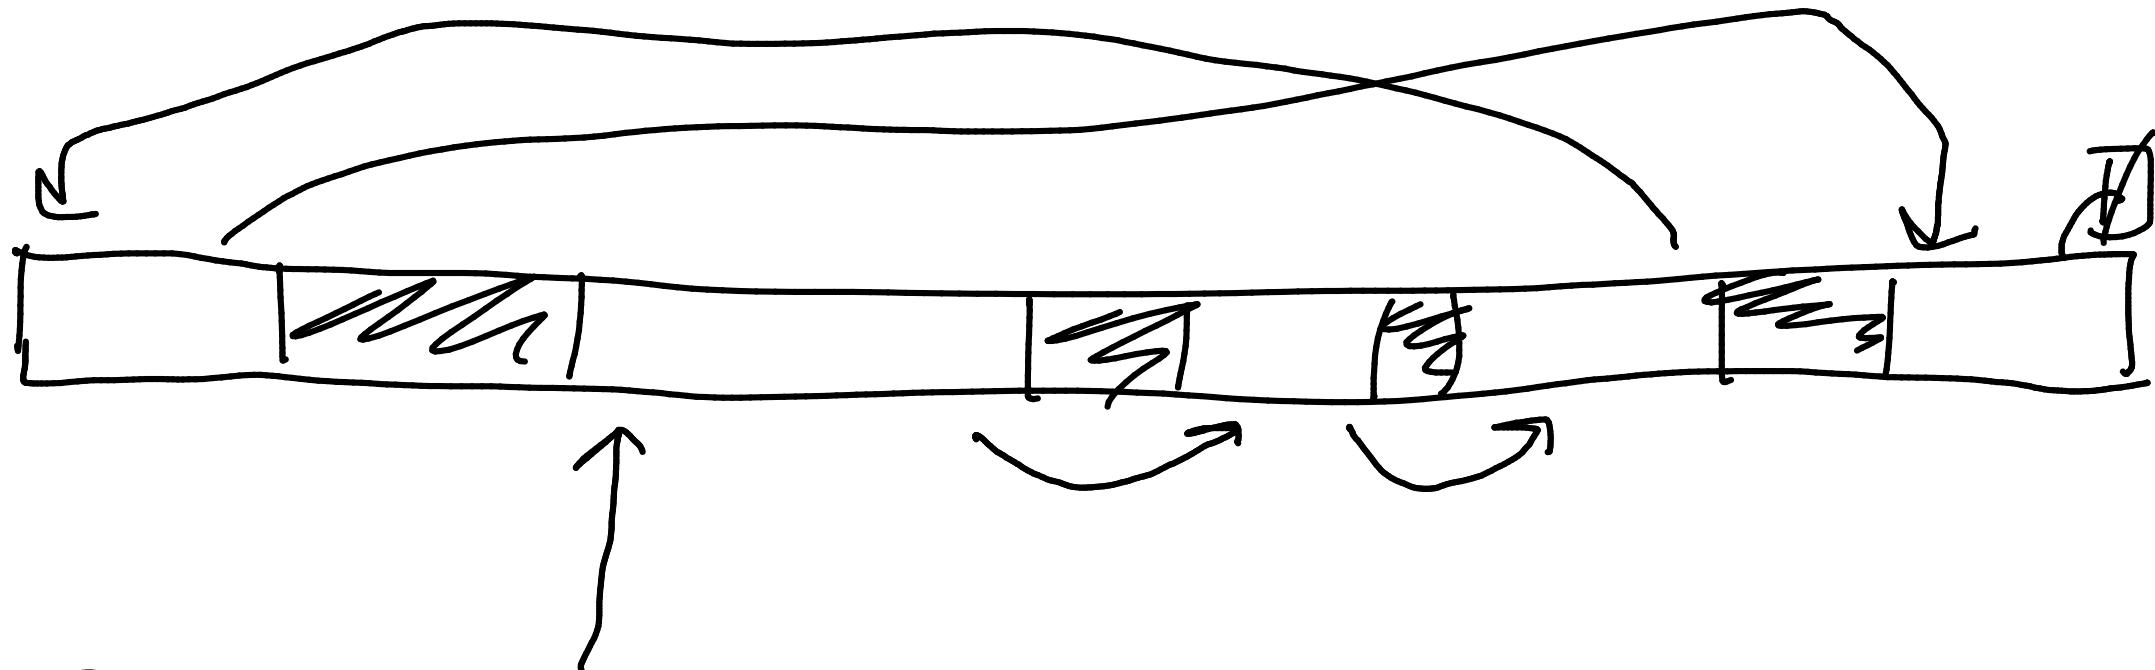
--- double --- messy ---

This is not double/allowable in
any language that exposes ptrs
in any meaningful way, like C
does

If you can't defrag memory,
how can you find / free / recombine
efficiently?

Lots of approaches!

If you did nothing: basic idea is you keep a list (a linked list) of free memory.



free memory

This is still just one linked list, which is slow to search for best fit.

Instead use lots of lists that represent various sizes of free memory

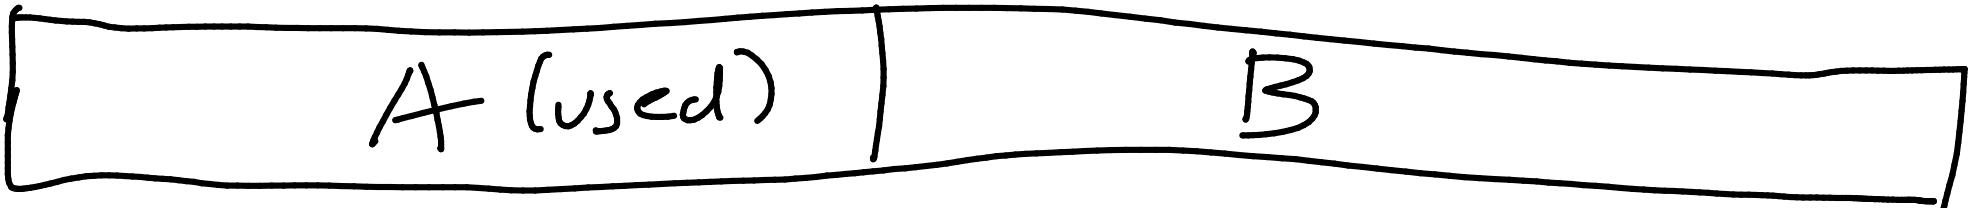
"segregated lists"

Buddy system - only break up
memory in sizes of
Powers of 2



← 64 bytes →

malloc(32)

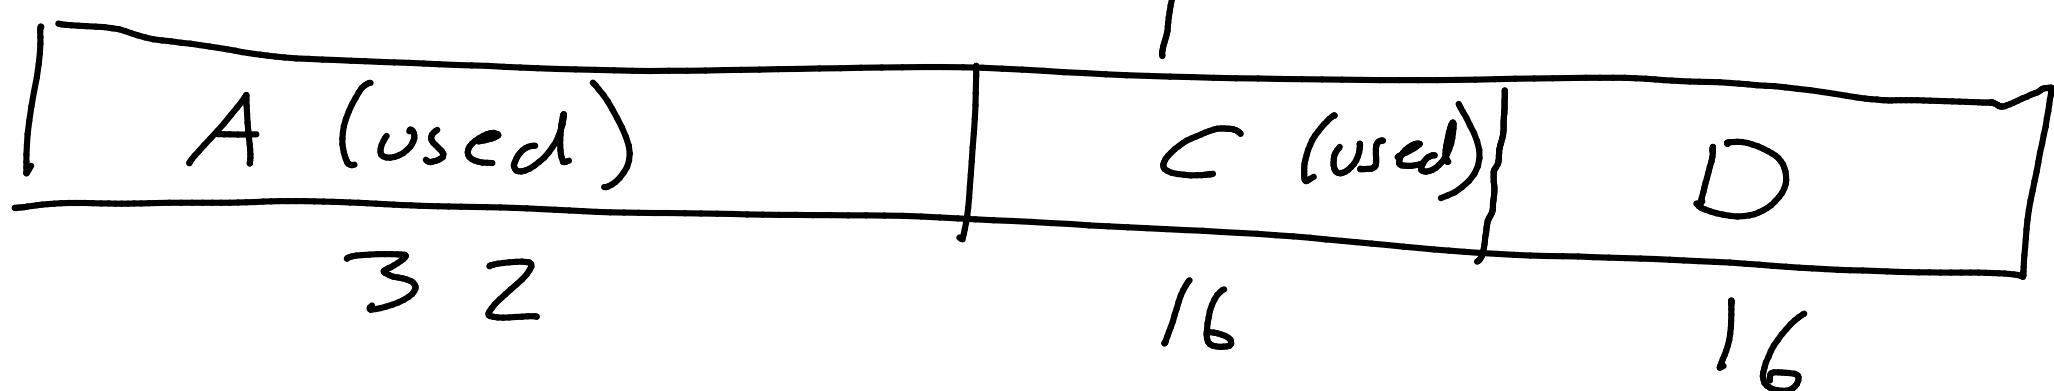


free lists

32: B

`malloc(15)` ↗ returned

Keep slicing memory until I have
a piece of size that's a power of 2,
but as small as possible

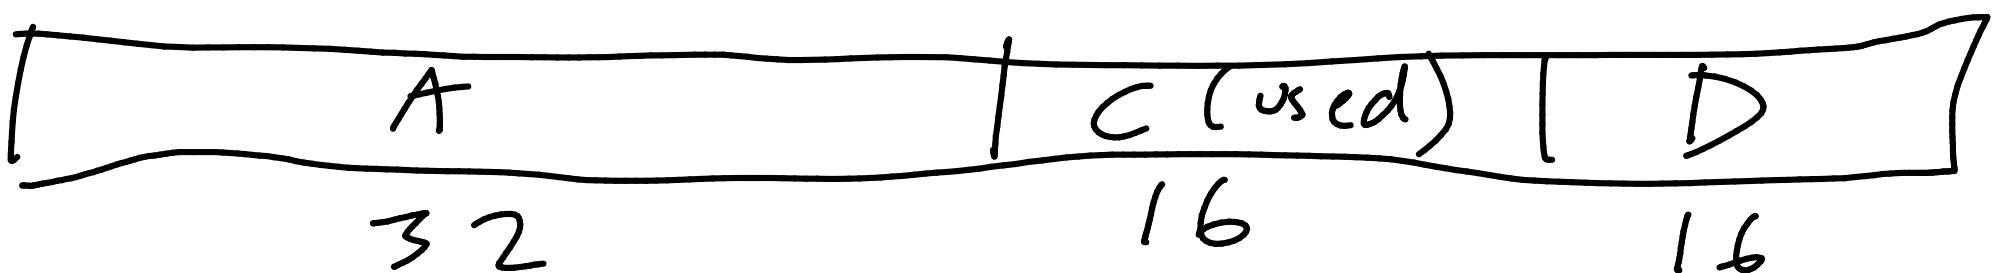


free lists

32: —

16: D

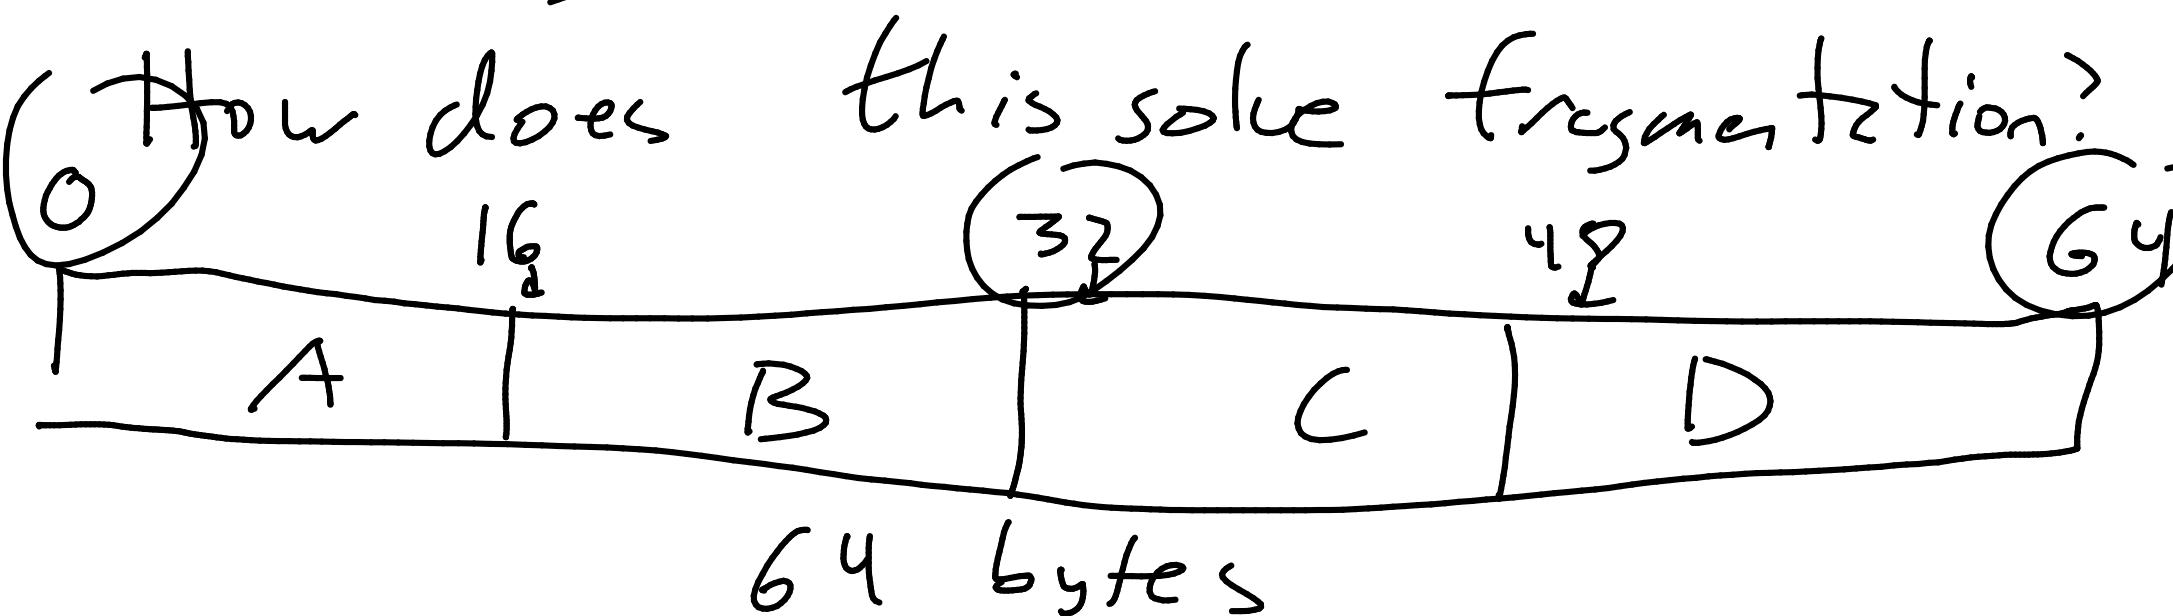
`free(A)` → A is unused, goes
back on the list



list:

32: A

16: D



malloc(16) \rightarrow A

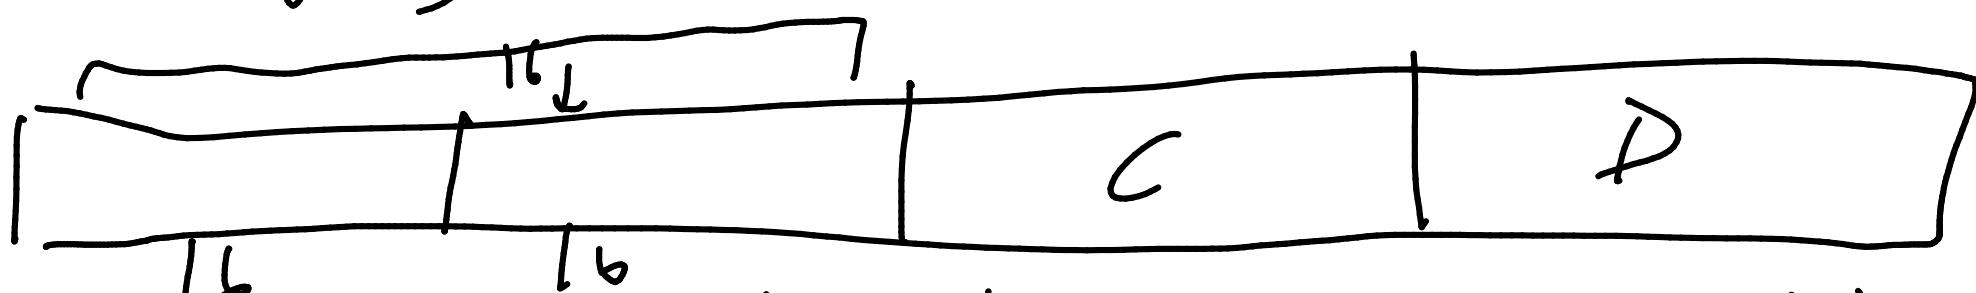
malloc(16) \rightarrow B

malloc(16) \rightarrow C

malloc(16) \rightarrow D

free(A)

free(B) 32



How does malloc know it can combine old A and B back together again?

Whenever anything is freed, see if its old buddy (the one it split from) is also free, and if so, recombine)

When I freed B, how do I know if its old buddy was to its left or its right?

B starts at 16

Chunks of size 32 must start at memory locs that are multiples of size 32

When I free B, look at mem address: if multiple of 32, it is a "left buddy", otherwise, it is a "right buddy"

But if memory you want to malloc
is not a power of 2, end up
wasting memory

Can we improve?

- Fibonacci numbers can help!
- seq where each number is the sum
of previous two

1, 1, 2, 3, 5, 8, 13, ...

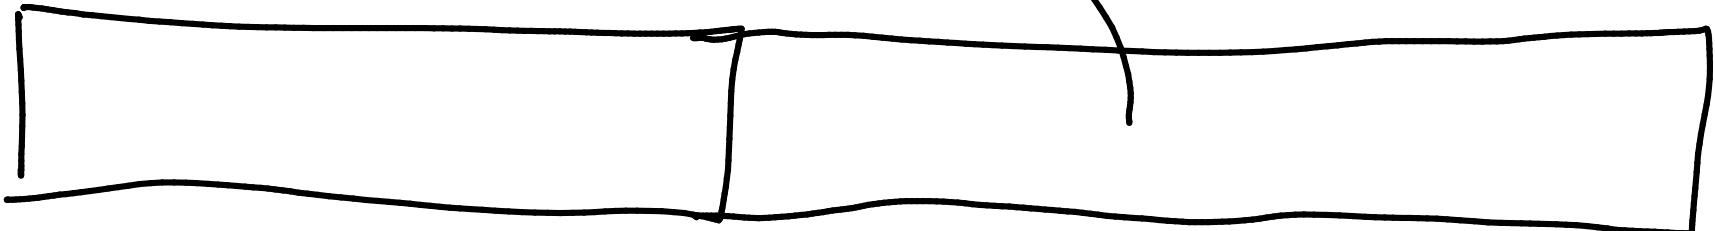
Grows more slowly than powers of 2

1 2 4 8 16 32 64

13



malloc(6)



5

8

Garbage collection

How could we have improved
talloc?

- talloc remembers all of *t*, but
never lets it go until you call *tfree*.

A good garbage collector would
free memory not in use anymore

Approach: reference counting
(easy)

Every time an object gets pointed to,
increase a count

When you move a ptr away,
decrease count

When count is 0, free it

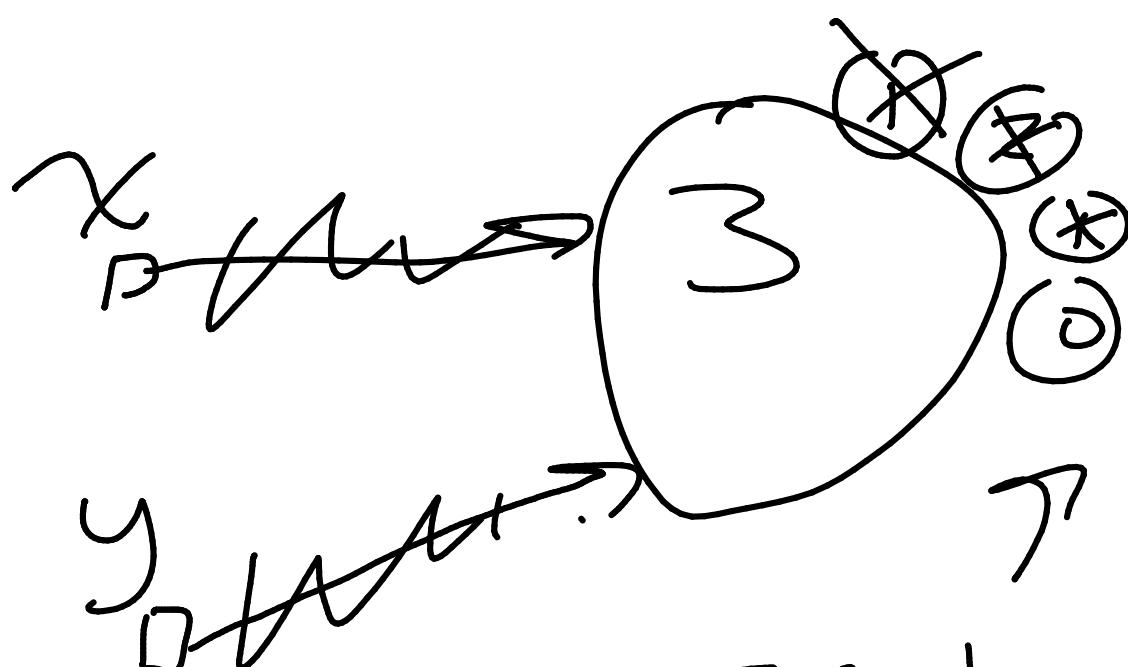
Python

$x = 3$

$y = x$

$y = \text{None}$

$x = \text{None}$



Zero!
Free the
memory

Advantages:

- easy to explain and to implement
- works incrementally
 - memory is freed right after you don't need it anymore

Amazing (except that it's broken)

Pseudocode

Node

Value

next

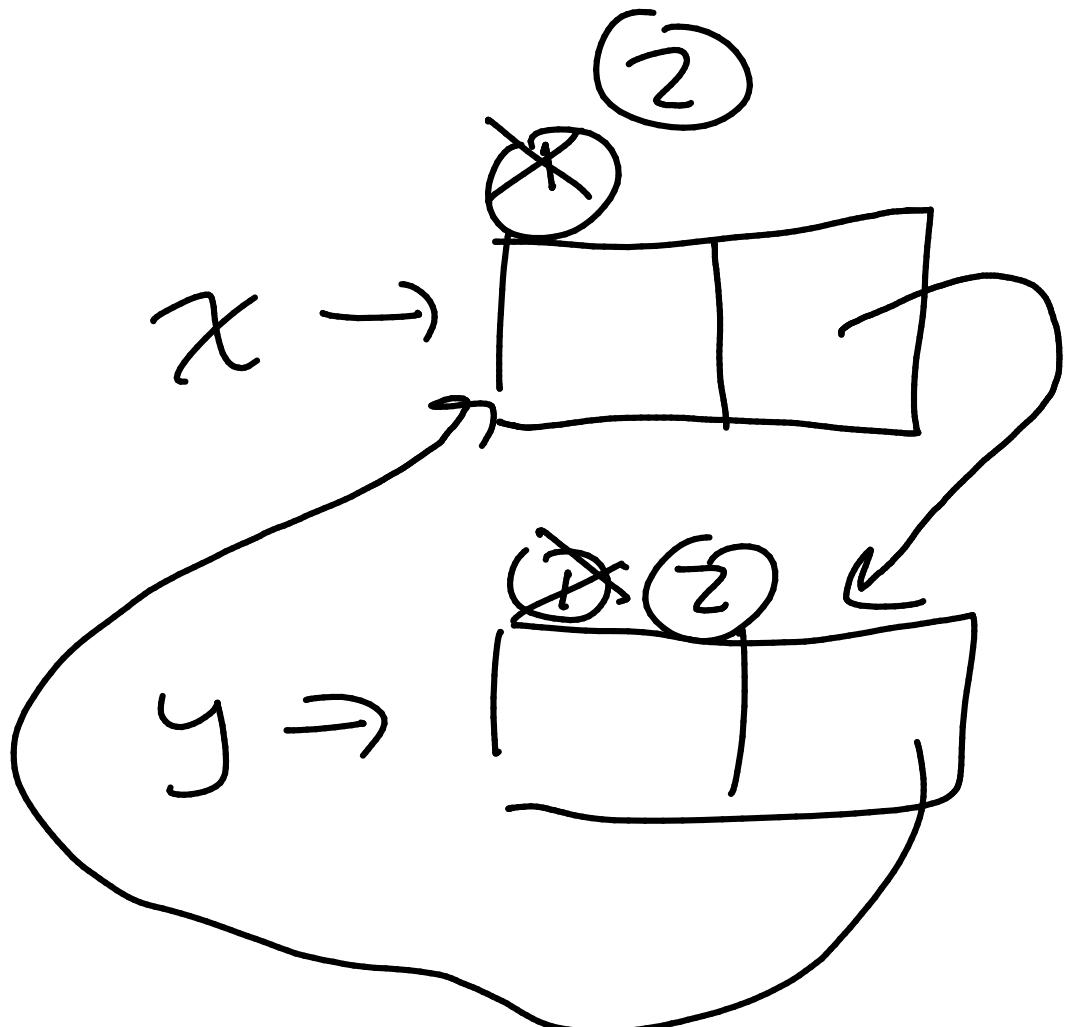
~~x = Node()~~

~~y = Node()~~

~~y = Node()~~

~~x.next = y~~

~~y.next = x~~



$x = \text{None}$

$y = \text{None}$

This fails if
have circular
references.

